

Lab 6 – Docker-swarm

Manage cluster of machines

We're going to build a cluster of Docker servers using Docker Swarm. You have 3 Docker servers:

- 1 manager
- 2 workers

Cluster initialization

1. Initialize the swarm cluster from the manager host using its IP from the same range than the 2 workers:

`docker swarm init --advertise-addr = 192.168.205.10`

```
vagrant@master:~$ docker swarm init --advertise-addr 192.168.205.10
Swarm initialized: current node (4osv85v370kxta3jheaojq4lu) is now a manager.
```

To add a worker to this swarm, run the following command:

```
docker swarm join --token SWMTKN-1-23t6832kuq0a2ctsjojwuwascafwwytrzbwk7ei
13qbyiy703d-4b8ijh7pbjg4389cs0jkhqo0z 192.168.205.10:2377
```

To add a manager to this swarm, run 'docker swarm join-token manager' and follow the instructions.

```
vagrant@master:~$
```

- What does this command generate?
- Check that the cluster has been created and that the manager is the single member.

```
vagrant@master:~$
vagrant@master:~$ docker node ls
```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS	ENGINE VER
4osv85v370kxta3jheaojq4lu	*	master	Ready	Active	Leader

```
23.0.1
vagrant@master:~$
```

- What is the token generated by this command used for? Display the 2 tokens.

```
vagrant@master:~$ docker swarm join-token manager
To add a manager to this swarm, run the following command:

    docker swarm join --token SWMTKN-1-23t6832kuq0a2ctsjojuwascafuytrzbwk7ei13qbyiy703d-85t82zag63pmpmqminquhj7 192.168.205.10:2377

vagrant@master:~$
vagrant@master:~$ docker swarm join-token worker
To add a worker to this swarm, run the following command:

    docker swarm join --token SWMTKN-1-23t6832kuq0a2ctsjojuwascafuytrzbwk7ei13qbyiy703d-4b8t1jh7pbjg4389cs0jkhqo0z 192.168.205.10:2377

vagrant@master:~$
```

2. Join worker1 and worker2 VMs to the cluster (as workers).

```
vagrant@worker1:~$ sudo docker swarm join --token SWMTKN-1-23t6832kuq0a2ctsjojuwascafuytrzbwk7ei13qbyiy703d-4b8t1jh7pbjg4389cs0jkhqo0z 192.168.205.10:2377
This node joined a swarm as a worker.
vagrant@worker1:~$
vagrant@worker1:~$ logout
brahim@Training:/sauvegarde/formations/formations_devops_TheTeam_Sofrecon_EnLigne/formation_devops_TheTeam_Sofrecon_EnLigne_21032023/Labs/docker-swarm/Lab6_docker-swarm$ vagrant ssh ff
sudo: Last login: Mon Mar 27 09:16:48 2023 from 10.0.2.2
vagrant@worker2:~$ sudo docker swarm join --token SWMTKN-1-23t6832kuq0a2ctsjojuwascafuytrzbwk7ei13qbyiy703d-4b8t1jh7pbjg4389cs0jkhqo0z 192.168.205.10:2377
This node joined a swarm as a worker.
vagrant@worker2:~$
vagrant@worker2:~$
```

- Check that the 2 workers are in the cluster and that they are ready.

```
vagrant@master:~$ docker node ls
```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS	ENGINE VER
40sv85v370kxta3jheaojq4lu *	master	Ready	Active	Leader	23.0.1
p77q5qex2w3yuzpa5byyuu0n9	worker1	Ready	Active		23.0.1
opbbzq8t0dq5njelx2kv086ki	worker2	Ready	Active		23.0.1

```
vagrant@master:~$
```

3. What networks are created following cluster initialization.

```
vagrant@master:~$ docker network ls
```

NETWORK ID	NAME	DRIVER	SCOPE
0306b44d2598	bridge	bridge	local
ea594837e6f0	docker_gwbridge	bridge	local
20030f024665	host	host	local
q852gchb4k4o	ingress	overlay	swarm
4ea2206efe1d	none	null	local

```
vagrant@master:~$
```

- Describe them.

```
vagrant@master:~$ docker inspect ingress
[
  {
    "Name": "ingress",
    "Id": "q852gchb4k4otvip80elj0xds",
    "Created": "2023-03-27T10:11:40.28792566Z",
    "Scope": "swarm",
    "Driver": "overlay",
    "EnableIPv6": false,
    "IPAM": {
      "Driver": "default",
      "Options": null,
      "Config": [
        {
          "Subnet": "10.0.0.0/24",
          "Gateway": "10.0.0.1"
        }
      ]
    },
    "Internal": false,
    "Attachable": false,
    "Ingress": true,
    "ConfigFrom": {
      "Network": ""
    },
    "ConfigOnly": false,
    "Containers": {
      "ingress-sbox": {
        "Name": "ingress-endpoint",
        "EndpointID": "00bf953ff05cc7587aa4d21c929deaeb50fda6ac013b567ca6f75e155e9c12",
        "MacAddress": "02:42:0a:00:00:02",

```

```
vagrant@master:~$ docker inspect docker_gwbridge
[
  {
    "Name": "docker_gwbridge",
    "Id": "ea594837e6f0889434fdb8ee14f919c6e940b65ec33a130ab66c09026bba3832",
    "Created": "2023-03-27T08:40:04.949902181Z",
    "Scope": "local",
    "Driver": "bridge",
    "EnableIPv6": false,
    "IPAM": {
      "Driver": "default",
      "Options": null,
      "Config": [
        {
          "Subnet": "172.18.0.0/16",
          "Gateway": "172.18.0.1"
        }
      ]
    },
    "Internal": false,
    "Attachable": false,
    "Ingress": false,
    "ConfigFrom": {
      "Network": ""
    },
    "ConfigOnly": false,
    "Containers": {
      "ingress-sbox": {
        "Name": "gateway_ingress-sbox",
        "EndpointID": "d9dde553a9e5df5b125269d0faedd526131c4b34cdca8a881f7ddb94f2bc6076",
        "IpAddress": "172.18.0.2/16",
        "IPv6Address": ""
      }
    }
  }
]
```

Launch of the Dockercoins application on Swarm

4. Create an overlay network to isolate all services of the application.

```
vagrant@master:~$ docker network create --driver overlay dockercoins-lan
j8h23tg4g96s9hi3numt1oz57
vagrant@master:~$ docker network ls
NETWORK ID        NAME                DRIVER            SCOPE
0306b44d2598      bridge              bridge            local
ea594837e6f0      docker_gwbridge     bridge            local
j8h23tg4g96s      dockercoins-lan     overlay           swarm
20030f024665      host                host              local
q852gchb4k40      ingress             overlay           swarm
4ea2206efe1d      none                null              local
vagrant@master:~$
```

- View details of our network

```
vagrant@master:~$ docker inspect dockercoins-lan
[
  {
    "Name": "dockercoins-lan",
    "Id": "j8h23tg4g96s9hi3numt1oz57",
    "Created": "2023-03-27T10:19:37.970030491Z",
    "Scope": "swarm",
    "Driver": "overlay",
    "EnableIPv6": false,
    "IPAM": {
      "Driver": "default",
      "Options": null,
      "Config": [
        {
          "Subnet": "10.0.1.0/24",
          "Gateway": "10.0.1.1"
        }
      ]
    },
    "Internal": false,
    "Attachable": false,
    "Ingress": false,
    "ConfigFrom": {
      "Network": ""
    },
    "ConfigOnly": false,
    "Containers": null,
    "Options": {
      "com.docker.network.driver.overlay.vxlanid_list": "4097"
    },
    "Labels": null
  }
]
```

5. Create the service running the redis database.

```
vagrant@master:~$ docker service create --name redis --network dockercoins-lan redis
u3dhpm79adoarmkod5coy735
overall progress: 1 out of 1 tasks
1/1: running
verify: Service converged
vagrant@master:~$
vagrant@master:~$ docker service ps redis
```

ID	NAME	IMAGE	NODE	DESIRED STATE	CURRENT STATE	ERROR	PORTS
hql2ep8x4wmg	redis.1	redis:latest	worker2	Running	Running 12 seconds ago		

```
vagrant@master:~$
```

- Wait until the server is in "running"
- Once redis is ready, run the hasher, rng, and worker services.

```
vagrant@master:~$ docker service create --name worker --network dockercoins-lan brahimhamdi/worker
2bv34v3fdauae3jy3e958wg5a
overall progress: 1 out of 1 tasks
1/1: running
verify: Service converged
vagrant@master:~$ docker service ps worker
```

ID	NAME	IMAGE	NODE	DESIRED STATE	CURRENT STATE	ERROR	PORTS
whgqd1ij4t2l	worker.1	brahimhamdi/worker:latest	master	Running	Running 11 seconds ago		

```
vagrant@master:~$
```

```
vagrant@master:~$ docker service create --name rng -p 8001:80 --network dockercoins-lan brahimhamdi/rng
1u8t47q9lpvdtir4dp0960f5w
overall progress: 1 out of 1 tasks
1/1: running [=====>]
verify: Service converged
vagrant@master:~$ docker service ps rng
```

ID	NAME	IMAGE	NODE	DESIRED STATE	CURRENT STATE	ERROR	PORTS
cdbwai02d6ru	rng.1	brahimhamdi/rng:latest	worker1	Running	Running 12 seconds ago		

```
vagrant@master:~$
```

```
vagrant@master:~$ docker service create --name hasher -p 8002:80 --network dockercoins-lan brahimhamdi/hasher
hsgdb4olyhiclk3x7rk2hn6c3
overall progress: 1 out of 1 tasks
1/1: running [=====>]
verify: Service converged
vagrant@master:~$ docker service ps hasher
```

ID	NAME	IMAGE	NODE	DESIRED STATE	CURRENT STATE	ERROR	PORTS
2pbma7qg223b	hasher.1	brahimhamdi/hasher:latest	master	Running	Running 11 seconds ago		

```
vagrant@master:~$
```

```
vagrant@master:~$ docker service create --name webui -p 8000:80 --network dockercoins-lan brahimhamdi/webui
hx0kisqujqe7i73jngcdkzquz
overall progress: 0 out of 1 tasks
overall progress: 1 out of 1 tasks
1/1: running [=====>]
verify: Service converged
vagrant@master:~$
vagrant@master:~$ docker service ps webui
```

ID	NAME	IMAGE	NODE	DESIRED STATE	CURRENT STATE	ERROR	PORTS
ydyasovxnsmo	webui.1	brahimhamdi/webui:latest	worker1	Running	Running 9 seconds ago		

```
vagrant@master:~$
```

- Check that these services are "running".

```
vagrant@master:~$ docker service ps redis worker rng hasher webui
```

ID	NAME	IMAGE	NODE	DESIRED STATE	CURRENT STATE
2pbma7qg223b	hasher.1	brahimhamdi/hasher:latest	master	Running	Running 6 minutes ago
hql2ep8x4wmg	redis.1	redis:latest	worker2	Running	Running 14 minutes ago
cdbwai02d6ru	rng.1	brahimhamdi/rng:latest	worker1	Running	Running 9 minutes ago
ydyasovxnsmo	webui.1	brahimhamdi/webui:latest	worker1	Running	Running 44 seconds ago
whgqd1ij4t2l	worker.1	brahimhamdi/worker:latest	master	Running	Running 12 minutes ago

```
vagrant@master:~$
```

- Give the IP address of each container as well as the node on which it is run.

Scaling on the swarm

6. scale out the worker service to 10 replicas.

```
vagrant@master:~$ docker service scale worker=10
worker scaled to 10
overall progress: 10 out of 10 tasks
1/10: running
2/10: running
3/10: running
4/10: running
5/10: running
6/10: running
7/10: running
8/10: running
9/10: running
10/10: running
verify: Service converged
```

- Once all the replicas are "running", display the webui. What is the rate of generation of dockercoins?



Current mining speed: ~31.9 hashes/second ([Tweet this!](#))

7. We want to run exactly one *rng* instance on each host. Unfortunately, this scheduling mode cannot be activated / deactivated for a existing service.

So, we have to start by removing the existing rng service, then recreate it with global scheduling


```
vagrant@master:~$ docker service create --mode global --name rng --network dockercoins-lan -p 8001:80 brahimhamdi/rng
n3j31o0c4a3gqey03v49da2kh
overall progress: 1 out of 3 tasks
overall progress: 3 out of 3 tasks
opbbzq8t0dq5: running
4osv85v370kx: running
p77q5qex2w3y: running
verify: Service converged
vagrant@master:~$
vagrant@master:~$ docker service ps rng
```


ID	NAME	IMAGE	NODE	DESIRED STATE	CURRENT STATE
PORTS					
py7zevxrgl8b	rng.4osv85v370kxta3jheaojq4lu	brahimhamdi/rng:latest	master	Running	Running about a minute ago
ps4x26we1to9	rng.opbbzq8t0dq5njlxl2kv086kl	brahimhamdi/rng:latest	worker2	Running	Running 10 seconds ago
w9bun6xip1z9	rng.p77q5qex2w3yuzpa5byyuu0n9	brahimhamdi/rng:latest	worker1	Running	Running 2 minutes ago

```
vagrant@master:~$
```

8. Check that only one instance of rng is running on each node.

```
ID          NAME          IMAGE          NODE          DESIRED STATE  CURRENT STATE
PORTS
py7zevxrgl8b  rng.4osv85v370kxta3jheaojq4lu  brahimhamdi/rng:latest  master        Running        Running about a minute ago
ps4x26we1to9  rng.opbbzq8t0dq5njlxl2kv086kl  brahimhamdi/rng:latest  worker2       Running        Running 10 seconds ago
w9bun6xip1z9  rng.p77q5qex2w3yuzpa5byyuu0n9  brahimhamdi/rng:latest  worker1       Running        Running 2 minutes ago
vagrant@master:~$
```

- Open the webui, what is the Dockercoins generation speed?



Current mining speed: ~24.0 hashes/second ([Tweet this!](#))

- Remove all services of the application.

```
vagrant@master:~$ docker service rm redis rng worker hasher webui
redis
rng
worker
hasher
webui
vagrant@master:~$ docker service ls
```

ID	NAME	MODE	REPLICAS	IMAGE	PORTS
----	------	------	----------	-------	-------

```
vagrant@master:~$
```

Docker stack

9. Now we will combine docker swarm and docker compose to be able to deploy on the cluster using docker-compose.yml file.

- Check that all services running

```
vagrant@master:~/lab4/dockercoins$ tree
```

```

.
├── docker-compose.yml
├── hasher
│   ├── Dockerfile
│   └── hasher.rb
├── rng
│   ├── Dockerfile
│   └── rng.py
├── webui
│   ├── Dockerfile
│   ├── Dockerfile.old
│   └── files
│       ├── d3.min.js
│       ├── index.html
│       ├── jquery-1.11.3.min.js
│       ├── jquery.js
│       ├── rickshaw.min.css
│       └── rickshaw.min.js
└── webui.js
└── worker
    ├── Dockerfile
    └── worker.py

```

```
5 directories, 16 files
```

```
vagrant@master:~/lab4/dockercoins$ docker stack deploy dockercoins -c docker-compose.yml
```

```
Creating network dockercoins_default
```

```
Creating service dockercoins_rng
```


```
Creating service dockercoins_hasher
```

```
Creating service dockercoins_webui
```

```
Creating service dockercoins_redis
```

```
Creating service dockercoins_worker
```

```
vagrant@master:~/lab4/dockercoins$
```



Current mining speed: ~4.0 hashes/second ([Tweet this!](#))